

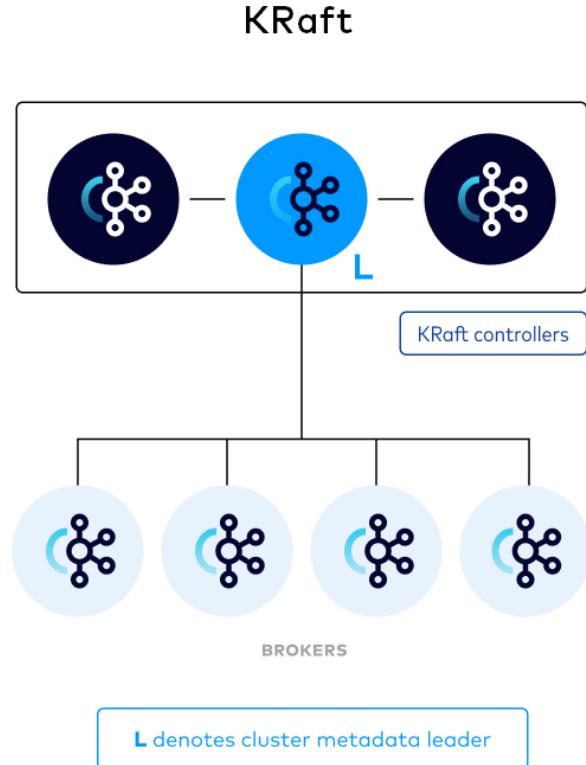
# Introduction to KRaft in Confluent Platform

## Simplifying Kafka Metadata Management

- **Kafka Raft (KRaft) Overview:** KRaft, short for Kafka Raft, is the consensus protocol that replaces ZooKeeper in Apache Kafka® starting from Confluent Platform 8.0. It consolidates metadata management directly into Kafka for improved simplicity and reliability.
- **Architectural Shift:** KRaft removes the external ZooKeeper dependency by embedding metadata management in Kafka itself. Each KRaft controller participates in a Raft quorum, ensuring strong consistency and streamlined control.
- **Motivation for Adoption:** The move to KRaft mode reduces operational complexity, enhances metadata propagation speed, and enables faster controller failover with fewer moving parts.
- **Key Benefit:** By eliminating ZooKeeper, Confluent simplifies deployment, increases fault tolerance, and supports larger Kafka clusters with lower maintenance overhead.

# KRaft Architecture

## How the Raft Quorum Manages Kafka Metadata



# KRaft Architecture

## How the Raft Quorum Manages Kafka Metadata



### Controller Quorum and Metadata Log

KRaft controllers form a Raft quorum that manages the Kafka metadata log, storing all topic, partition, ISR, and configuration changes. This ensures consistency and durability across controllers.



### Active and Follower Controllers

The active controller handles all broker RPCs, while follower controllers replicate metadata updates and serve as hot standbys, ensuring immediate recovery during failover.



### Leader Election and Consensus

Using the Raft protocol, KRaft achieves consensus without external coordination. Majority voting enables leader election, fault tolerance, and fast recovery.



### Snapshot and Recovery Mechanism

Controllers periodically write metadata snapshots to disk, improving recovery times and reducing replay overhead during restarts.

# Scaling Kafka with KRaft

## Performance, Partition Limits, and Failover Improvements

- **Partition and Cluster Limits:** Kafka scalability in KRaft mode depends on both the per-node partition count limit and the cluster-wide partition limit. Adding more nodes remains the primary way to increase capacity.
- **Enhanced Metadata Propagation:** KRaft's quorum controller minimizes metadata propagation delay, enabling faster synchronization across brokers and reducing cluster downtime.
- **Rapid Failover:** Experiments by Confluent demonstrate near-instantaneous controller failover in KRaft. In a 2-million-partition cluster, KRaft achieved 10× the partition scale of ZooKeeper-based clusters.
- **High Availability through Majority Consensus:** A three-node quorum can survive one controller failure; five nodes can survive two. Majority voting ensures operational continuity during failures.

# KRaft Configuration Essentials

## Core Settings Required for Confluent Platform Deployment

- **process.roles**: Specifies whether the Kafka server acts as a controller, broker, or both (though combined mode is not supported for production). This setting activates KRaft mode.
- **node.id**: Assigns a unique identifier to each node. No two servers—controller or broker—may share the same ID within a cluster.
- **controller.quorum.voters**: Defines the set of controllers participating in the Raft quorum using the format {id}@{host}:{port}. Ensures all nodes are aware of the quorum members.
- **controller.listener.names**: Lists controller listeners and ensures they match the listeners property. Essential for isolated or combined mode configurations.

# Listeners, Logs, and Security in KRaft Mode

## Ensuring Secure and Reliable Metadata Communication

- **Inter-Broker Communication:** Defined via either `inter.broker.listener.name` or `security.inter.broker.protocol`. These settings ensure secure communication between brokers and controllers.
- **Listener Configuration:** Controllers must list their own listeners under `controller.listener.names`, consistent with the quorum voters list. Brokers and controllers use mapped security protocols for communication.
- **Metadata Log Directories:** The `metadata.log.dir` property specifies where controller metadata is stored. If unset, it defaults to the first directory in `log.dirs`.
- **Security and Truststores:** Secure communication is achieved by defining SSL keystores and truststores per listener. Controllers must include truststore configurations for broker communication.

# Tools for Managing and Debugging KRaft Mode

Command-Line Utilities for Cluster Insight and Diagnostics



## kafka-metadata-quorum

Describes the runtime status of the KRaft metadata partition. Can connect via `--bootstrap-server` or `--bootstrap-controller` to inspect quorum state and leadership.



## kafka-dump-log

Decodes and prints metadata records from KRaft log segments and snapshots, enabling deep inspection of stored metadata.



## kafka-metadata-shell

Provides an interactive shell for exploring metadata partitions. Supports secure connections to running controllers via SSL.



## kafka-migration-check

Assesses migration status and readiness for KRaft clusters. Useful for validating Confluent Platform transitions post-ZooKeeper.

# Migration, Limitations & Best Practices

## Transitioning from ZooKeeper and Ensuring KRaft Stability



### Migration Process

Use the `kafka-migration-check` tool to validate migration readiness. Ensure each node's storage is formatted with a consistent cluster ID using the `kafka-storage format` command.



### Unsupported Combined Mode

Running a Kafka node as both broker and controller (combined mode) is not supported in production due to security and operational gaps.



### Quorum Reconfiguration

Prior to Confluent Platform 7.9, KRaft clusters could not add or remove controllers dynamically. Later versions introduce controlled quorum reconfiguration.



### Monitoring Limitations

Certain tools, such as Health+, may still misreport KRaft controllers as brokers, impacting alert accuracy. Administrators should account for this discrepancy.

# Conclusion: The Benefits of KRaft Architecture

Simpler, Faster, and More Reliable Kafka Metadata Management

- **Unified Metadata Management:** KRaft embeds metadata handling directly into Kafka, eliminating ZooKeeper and reducing operational overhead.
- **Improved Fault Tolerance:** Raft-based consensus and majority voting ensure resilient controller failover and minimal downtime.
- **Operational Efficiency:** Simplified configuration and deployment enable easier scaling and reduced administrative complexity.
- **Enhanced Scalability:** Supports up to 10× more partitions than ZooKeeper-based clusters, allowing massive deployments to operate efficiently.