

Monitoring Confluent Kafka

Comprehensive Overview



Purpose of Monitoring

Ensures Kafka cluster health, availability, and performance by tracking key metrics such as throughput, latency, and replication status.



Confluent Integration

Confluent Platform extends Apache Kafka monitoring with Control Center, Health+, and monitoring interceptors for end-to-end visibility.



Multi-Tool Ecosystem

Supports JMX-based metrics, Prometheus & Grafana integrations, and REST APIs for observability and alerting across distributed environments.



Objective of This Presentation

Provide a structured overview of Kafka monitoring concepts, metrics, and tools to enable proactive operations management.

Importance of Kafka Monitoring

Ensuring Reliability, Performance, and Availability

- **Operational Stability:** Continuous monitoring prevents outages by detecting broker failures, controller transitions, and partition unavailability early.
- **Performance Optimization:** Throughput, latency, and resource utilization metrics help fine-tune cluster performance and resource allocation.
- **Data Reliability:** Replication, ISR status, and unclean leader election metrics ensure data integrity and fault tolerance across brokers.
- **Proactive Troubleshooting:** Monitoring allows early detection of lag, bottlenecks, and misconfigurations before they impact production workloads.

Monitoring Architecture Overview

Data Flow and Key Components

- **JMX Metrics Foundation:** Kafka exposes internal metrics via JMX, enabling collection of broker and client-level performance data for external monitoring tools.
- **Confluent Control Center:** Provides a centralized web interface for monitoring brokers, topics, producers, and consumers across the Confluent Platform.
- **Prometheus & Grafana Integration:** JMX Exporter and Prometheus agents collect and visualize Kafka metrics, allowing historical trend analysis and alerting.
- **Health+ Cloud Telemetry:** Offers intelligent, cloud-based alerts and dashboards derived from best practices across thousands of Kafka clusters.

JMX Monitoring Foundation

Core Metric Interface in Kafka

- **Role of JMX in Kafka:** JMX (Java Management Extensions) serves as the primary mechanism for exposing broker and client metrics, enabling real-time visibility into Kafka operations.
- **Configuration Essentials:** Set JMX_PORT, JMX_HOSTNAME, and KAFKA_JMX_OPTS when starting brokers. Secure production setups with authentication and SSL to prevent unauthorized access.
- **Metric Exposure and Visualization:** Metrics are emitted via Yammer (brokers) and Kafka Metrics (clients) and can be visualized with tools like JConsole or exported to Prometheus.
- **Operational Significance:** Provides granular observability across threads, queues, and topics – essential for diagnosing latency, throughput, and replication anomalies.

Confluent Monitoring Interceptors

Client-Side Metrics for Stream Insight



Purpose and Function

Lightweight client libraries that collect real-time statistics on messages sent and received, feeding data into Confluent Control Center.



Key Configuration Properties

Includes `confluent.monitoring.interceptor.topic`, `publishMs`, and `client.id` – defining topic destination, interval, and logical client identity.



Integration Simplicity

Easily embedded into producers and consumers with minimal performance overhead and no code changes.



Operational Benefits

Enables per-application and per-topic monitoring, enhancing visibility across distributed microservices using Kafka.

Critical Broker Metrics

Key Indicators of Kafka Cluster Health

- **ActiveControllerCount:** Indicates the number of active controllers in the cluster; must equal 1 across all brokers. Deviations signal leadership issues.
- **OfflinePartitionsCount:** Represents partitions without an active leader, rendering them unavailable for reads and writes. Any nonzero value requires immediate attention.
- **UncleanLeaderElectionsPerSec:** Tracks rate of unsafe leader elections that risk data loss. Should always remain at 0 to maintain consistency.
- **Alerting and Remediation:** Set critical alerts for these metrics; ensure automated failover handling and replication verification mechanisms are in place.

Replication and ISR Metrics

Monitoring Data Consistency and Cluster Resilience



Understanding ISR

In-Sync Replicas (ISR) are replicas fully caught up with the partition leader. A stable ISR set ensures data durability and high availability.



Key ISR Parameters

`min.insync.replicas` defines minimum acknowledgments for writes; `replica.lag.time.max.ms` determines allowable lag before replica removal from ISR.



Critical ISR Metrics

Monitor `IsrShrinksPerSec` and `IsrExpandsPerSec` to detect broker failures and recoveries; `InSyncReplicasCount` and `UnderMinIsr` for partition health.



Operational Insight

Frequent ISR shrink or expansion indicates instability. Establish baselines and correlate with broker health and network latency metrics.

Consumer Lag Monitoring

Tracking Message Consumption Efficiency

- **Definition and Importance:** Consumer lag measures how far behind a consumer group is from the latest offset – a direct indicator of processing delays.
- **Manual Monitoring:** Use `kafka-consumer-groups.sh` to list, describe, and track consumer lag through `CURRENT-OFFSET`, `LOG-END-OFFSET`, and `LAG` columns.
- **Automated Monitoring:** Integrate Prometheus, Grafana, or LinkedIn Burrow to visualize lag trends and automate threshold-based alerts.
- **Confluent Enhancements:** Enable `confluent.consumer.lag.emitter` for JMX-based lag metrics and real-time tracking in Confluent Control Center.

Producer and Consumer Metrics

Measuring Throughput, Latency, and Efficiency



Producer Performance Metrics

Key metrics include record-send-rate, record-error-rate, request-latency-avg, and batch-size-avg — essential for evaluating data ingestion efficiency.



Consumer Fetch Metrics

bytes-consumed-rate, records-consumed-rate, and fetch-latency-avg reveal how quickly consumers retrieve and process data.



Lag and Throttling Indicators

records-lag-max and fetch-throttle-time-avg help identify processing delays and quota-enforced slowdowns.



Operational Optimization

Monitoring producer and consumer metrics together enables balancing of throughput, tuning batch sizes, and diagnosing backpressure.

Network and Disk Metrics

Ensuring Throughput and Storage Reliability



Network Request Metrics

Monitor RequestQueueTimeMs, LocalTimeMs, RemoteTimeMs, and TotalTimeMs to analyze request handling and end-to-end latency.



Queue Health Indicators

Track RequestQueueSize and ResponseQueueSize for congestion. High values indicate potential processing bottlenecks.



Log Retention Metrics

LogEndOffset, LogStartOffset, and Size metrics ensure message retention policies function correctly within storage limits.



Disk Utilization

Monitor disk usage and replication factors to prevent space exhaustion and ensure healthy log segment management.

Confluent Control Center

Centralized Monitoring and Management Platform

- **Unified Kafka Monitoring:** Provides a real-time overview of cluster health, broker metrics, consumer lag, and topic throughput across the Confluent ecosystem.
- **Operational Features:** Supports topic management, ksqlDB query monitoring, and Connect/Replicator oversight from a single web interface.
- **Alerting Capabilities:** Configurable alerts for anomalies such as consumer lag, cluster downtime, and partition unavailability.
- **Access and Integration:** Accessible at <http://localhost:9021>, integrates seamlessly with brokers, Connect, and Control Center telemetry collectors.

Confluent Health+

Cloud-Based Kafka Monitoring and Proactive Insights

- **Intelligent Alerts:** Provides over 50 pre-tuned alerts derived from operating 5,000+ clusters, covering anomalies in brokers, consumers, and connectors.
- **Monitoring Dashboards:** Offers a centralized, cloud-hosted interface with historical data visualization and trend analysis for all connected clusters.
- **Seamless Integration:** Exports metrics to Prometheus, Slack, Microsoft Teams, or email for unified alerting and collaboration.
- **Automated Telemetry Setup:** Enable Health+ by configuring Telemetry Reporter in Kafka server properties or dynamically using kafka-configs.

Prometheus & Grafana Integration

Open-Source Observability for Kafka

- **JMX Exporter Setup:** Attach the JMX Prometheus Java agent to brokers using the KAFKA_OPTS variable to expose metrics over HTTP for Prometheus scraping.
- **Prometheus Collection:** Collects and stores Kafka metrics as time-series data, enabling query-based monitoring and alerting using PromQL.
- **Grafana Visualization:** Dashboards display real-time trends in producer/consumer throughput, request latency, and broker health.
- **Example Metrics:** Monitor request rate, producer errors, and latency percentiles (e.g., P99) to detect anomalies and capacity issues.

Kafka Connect, ksqlDB, and Schema Registry Monitoring

Extending Observability Across Kafka Ecosystem Services

- **Kafka Connect Metrics:** Monitor task-count, connector-count, and startup success via JMX (kafka.connect:type=connect-worker-metrics) and REST API endpoints.
- **ksqldb Insights:** Track running-queries, error-queries, and message throughput using JMX metrics for ksql-engine-query-stats to optimize query performance.
- **Schema Registry Metrics:** Measure active connections, request-latency-avg, and error rates to ensure schema validation service reliability and response speed.
- **Unified Observability:** Integrate these metrics into Control Center or Prometheus for cross-component visibility and proactive troubleshooting.

ZooKeeper Monitoring (Legacy Mode)

Ensuring Coordination Service Health



Role in Kafka

ZooKeeper manages metadata, leader election, and configuration synchronization for legacy Kafka clusters not running in KRaft mode.



Key Metrics

Monitor outstanding-requests, avg-latency, num-alive-connections, followers, and pending-syncs to assess service responsiveness.



Alert Thresholds

Trigger alerts for avg-latency > 100ms, outstanding requests growing consistently, or follower counts dropping below quorum.



Transition to KRaft

Modern deployments are migrating to KRaft mode, reducing dependency on ZooKeeper while improving scalability and manageability.

Alerting Best Practices

Proactive Response and Reliability Safeguards



Critical Thresholds

Define strict thresholds for metrics such as ActiveControllerCount ≠ 1, OfflinePartitionsCount > 0, and Consumer Lag > baseline × 2.



Multi-Channel Notifications

Deliver alerts via email, Slack, SMS, or PagerDuty for rapid response and escalation management.



SLO-Driven Monitoring

Align alert configurations with service-level objectives (SLOs) and SLAs to ensure operational consistency.



Continuous Optimization

Regularly review alert trends, tune thresholds, and perform capacity planning to prevent alert fatigue and enhance reliability.

Logging Configuration

Managing Kafka Logs for Observability and Debugging

- **Log4j Configuration:** Define logging behavior using KAFKA_LOG4J_OPTS to specify log4j.properties for brokers, Connect, and other components.
- **Dynamic Log Level Changes:** Adjust log levels in real-time using kafka-configs.sh for brokers or REST API for Kafka Connect without restarting services.
- **Granular Logging Control:** Supports per-component and per-package log level tuning (INFO, DEBUG, TRACE) for targeted diagnostics.
- **Best Practices:** Maintain separate log files for system, request, and GC logs. Regularly rotate logs and monitor disk usage to prevent overflow.

Monitoring Tools Comparison

Choosing the Right Solution for Your Environment

- **Confluent Control Center:** Enterprise-grade native Kafka monitoring with UI-based topic, consumer, and alert management.
- **Health+:** Cloud-based SaaS for proactive monitoring and intelligent alerting across multiple Kafka clusters.
- **Prometheus + Grafana:** Open-source stack offering customizable dashboards, historical analysis, and flexible alerting via PromQL.
- **Third-Party Tools:** Datadog provides full-stack observability; LinkedIn Burrow focuses on consumer lag; ELK Stack supports log aggregation and search.

Summary of Key Monitoring Areas

Integrating Metrics for End-to-End Kafka Observability



Broker Health

Track controller status, partition availability, and replication metrics to ensure stability and high availability.



Consumer Performance

Monitor lag, fetch latency, and consumption rate to confirm that message processing matches production pace.



Producer Efficiency

Assess send rate, error rate, and request latency to maintain high data ingestion performance and reliability.



Resource Utilization

Measure CPU, memory, disk I/O, and network metrics to anticipate bottlenecks and perform capacity planning.

Conclusion

Building a Resilient Kafka Monitoring Ecosystem



Holistic Monitoring

Effective observability integrates metrics, logs, and alerts across brokers, producers, consumers, and supporting services.



Proactive Management

Automated alerting and Health+ insights help detect anomalies early and prevent downtime in production systems.



Toolchain Synergy

Combine Control Center's enterprise UI with Prometheus and Grafana's open-source flexibility for comprehensive coverage.



Sustainable Reliability

Continuous review, trend analysis, and performance tuning strengthen Kafka's resilience and scalability.