

## **Exercise:** monitoring stack for **Confluent Kafka** using **Prometheus** (for metric collection) and **Grafana** (for visualization).

### **Objective**

By the end of this exercise, you will have:

1. A running Confluent Kafka Cluster (Single Broker).
  2. A JMX Exporter attached to Kafka to expose metrics.
  3. Prometheus scraping the metrics.
  4. A Grafana Dashboard visualizing real-time throughput, lag, and broker health.
- 

### **Prerequisites**

- **Docker** and **Docker Compose** installed on your machine.
  - Basic familiarity with command-line operations.
  - Internet access (to download Docker images and the JMX agent jar).
- 

To install Docker and Docker Compose on Ubuntu, follow these steps:

1. Update System Packages:

```
sudo apt update  
sudo apt upgrade -y
```

2. Install Prerequisite Packages:

```
sudo apt install -y ca-certificates curl gnupg lsb-release
```

3. Add Docker's Official GPG Key:

```
sudo mkdir -p /etc/apt/keyrings  
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o  
/etc/apt/keyrings/docker.gpg
```

4. Set up the Docker Repository:

```
echo \  
"deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]  
https://download.docker.com/linux/ubuntu \  
$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

5. Install Docker Engine and Docker Compose Plugin:

```
sudo apt update
```

```
sudo apt install -y docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-plugin
```

6. Verify Docker Installation:

```
sudo docker run hello-world
```

This command should download a test image and run a container, printing a "Hello from Docker!" message if successful.

7. Add Your User to the docker Group (Optional, but Recommended):

To run Docker commands without sudo, add your user to the docker group:

```
sudo usermod -aG docker ${USER}
```

You will need to log out and log back in for this change to take effect.

8. Verify Docker Compose Installation:

```
docker compose version
```

This command should display the installed version of Docker Compose.

---

## Step 1: Project Workspace Setup

Create a directory for your project and the necessary subfolders to hold configurations.

```
mkdir kafka-monitoring-lab
```

```
cd kafka-monitoring-lab
```

```
mkdir -p config/prometheus jars
```

---

## Step 2: Download JMX Prometheus Agent

The JMX Exporter acts as a "bridge" between Kafka's internal Java metrics (JMX) and Prometheus. We need to download the Java agent .jar file.

1. Download the jar into your jars folder:

```
curl -o jars/jmx_prometheus_javaagent.jar \
```

```
https://repo1.maven.org/maven2/io/prometheus/jmx/jmx_prometheus_javaagent/0.19.0/jmx_prometheus_javaagent-0.19.0.jar
```

---

### Step 3: Create Configuration Files

#### A. JMX Exporter Config (config/kafka\_jmx\_config.yml)

Create this file to tell the agent which metrics to export and how to format them.

**File:** kafka-monitoring-lab/config/kafka\_jmx\_config.yml

lowercaseOutputName: true

lowercaseOutputLabelNames: true

rules:

# Broker Topic Metrics (BytesIn, BytesOut, MessagesIn, etc.)

- pattern: kafka.server<type=BrokerTopicMetrics, name=(.+), topic=(.+)><>Count

name: kafka\_server\_brokertopicmetrics\_\$1\_total

type: COUNTER

labels:

topic: "\$2"

- pattern: kafka.server<type=BrokerTopicMetrics, name=(.+)><>Count

name: kafka\_server\_brokertopicmetrics\_\$1\_total

type: COUNTER

# Request Handling (Produce, Fetch, etc.)

- pattern: kafka.network<type=RequestMetrics, name=RequestsPerSec, request=(.+)><>Count

name: kafka\_network\_requestmetrics\_requests\_total

type: COUNTER

labels:

request: "\$1"

# JVM Metrics (Memory, GC, Threads)

- pattern: "java.lang<type=Memory><HeapMemoryUsage>used"

```
name: jvm_memory_bytes_used
type: GAUGE
- pattern: "java.lang<type=GarbageCollector, name=(.+)><CollectionCount>Count"
  name: jvm_gc_collection_seconds_count
  type: COUNTER
  labels:
    gc: "$1"
```

## B. Prometheus Config (config/prometheus.yml)

Create this file to tell Prometheus where to find the Kafka metrics.

**File:** kafka-monitoring-lab/config/prometheus.yml

global:

```
scrape_interval: 10s
```

scrape\_configs:

```
- job_name: 'kafka-broker'
```

static\_configs:

```
- targets: ['kafka:7071']
```

labels:

```
env: 'dev'
```

---

## Step 4: Create Docker Compose File

This file defines our entire stack. We are injecting the KAFKA\_OPTS environment variable to attach the JMX agent to the Kafka process.

**File:** kafka-monitoring-lab/docker-compose.yml

```
version: '3.8'
```

services:

zookeeper:

```
image: confluentinc/cp-zookeeper:7.5.0
```

```
hostname: zookeeper
container_name: zookeeper
environment:
  ZOOKEEPER_CLIENT_PORT: 2181
  ZOOKEEPER_TICK_TIME: 2000

kafka:
  image: confluentinc/cp-server:7.5.0
  hostname: kafka
  container_name: kafka
  depends_on:
    - zookeeper
  ports:
    - "9092:9092"
    - "7071:7071" # Exposing JMX Exporter port
  volumes:
    - ./jars:/usr/share/jmx_exporter/
    - ./config/kafka_jmx_config.yml:/usr/share/jmx_exporter/config.yml
  environment:
    KAFKA_BROKER_ID: 1
    KAFKA_ZOOKEEPER_CONNECT: 'zookeeper:2181'
    KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
      PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
    KAFKA_ADVERTISED_LISTENERS:
      PLAINTEXT://kafka:29092,PLAINTEXT_HOST://localhost:9092
    KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
    KAFKA_GROUP_INITIAL_REBALANCE_DELAY_MS: 0
    KAFKA_CONFLUENT_LICENSE_TOPIC_REPLICATION_FACTOR: 1
    KAFKA_CONFLUENT_BALANCER_TOPIC_REPLICATION_FACTOR: 1
```

```
# CRITICAL: Attach the Java Agent here  
KAFKA_OPTS: "-  
javaagent:/usr/share/jmx_exporter/jmx_prometheus_javaagent.jar=7071:/usr/share/jmx_ex  
porter/config.yml"
```

prometheus:

```
image: prom/prometheus  
container_name: prometheus
```

ports:

```
- "9090:9090"
```

volumes:

```
- ./config/prometheus.yml:/etc/prometheus/prometheus.yml
```

command:

```
- '--config.file=/etc/prometheus/prometheus.yml'
```

grafana:

```
image: grafana/grafana  
container_name: grafana
```

ports:

```
- "3000:3000"
```

environment:

```
- GF_SECURITY_ADMIN_PASSWORD=admin
```

depends\_on:

```
- prometheus
```

---

## Step 5: Launch the Stack

1. Start the services:

```
docker-compose up -d
```

2. Verify everything is running:

```
docker-compose ps
```

*(Ensure kafka, zookeeper, prometheus, and grafana are all "Up")*

**3. Verify Metrics Export:** Open your browser to <http://localhost:7071>.

- You should see a long page of text metrics. If you see this, the JMX Exporter is working!

**4. Verify Prometheus Scrape:** Open <http://localhost:9090/targets>.

- The state for kafka-broker should be **UP** (green).
- 

## **Step 6: Generate Kafka Traffic**

To see interesting charts, we need data moving through the system.

**1. Enter the Kafka container:**

```
docker exec -it kafka /bin/bash
```

**2. Create a topic:**

```
kafka-topics --bootstrap-server localhost:9092 --create --topic test-topic --partitions 3 --replication-factor 1
```

**3. Run the performance test producer (generates load):**

```
kafka-producer-perf-test \  
--topic test-topic \  
--num-records 500000 \  
--record-size 1000 \  
--throughput 5000 \  
--producer-props bootstrap.servers=localhost:9092
```

*(Let this run in the background or open a new terminal tab to keep it running)*

---

## **Step 7: Configure Grafana**

**1. Login:**

- Go to <http://localhost:3000>
- User: admin, Password: admin (Skip password change if asked)

**2. Add Data Source:**

- Go to **Connections (or Configuration) > Data Sources > Add data source.**
- Select **Prometheus**.
- Set URL to: <http://prometheus:9090> (Note: we use the container name, not localhost).
- Click **Save & Test**. You should see "Data source is working".

### 3. Import Dashboard:

- Go to **Dashboards > New > Import**.
  - In "Import via grafana.com", enter ID: **721** (A popular Kafka Overview dashboard) or **18276**.
  - Click **Load**.
  - Select your Prometheus data source in the dropdown.
  - Click **Import**.
- 

## Step 8: Analyze the Results

Look at your new Dashboard. You should now see:

- **Incoming Message Rate**: Spiking up to ~5k ops/sec (matches your perf-test).
- **Bytes In/Out**: Showing the network throughput.
- **Active Controller Count**: Should be 1.
- **Offline Partitions**: Should be 0.

## Troubleshooting

- **Kafka container exits immediately**: Usually a memory issue. Increase Docker memory limit to at least 4GB.
- **Target DOWN in Prometheus**: Check <http://localhost:7071>. If that fails, check docker logs kafka. Look for "javaagent" errors.
- **No Data in Grafana**: Ensure the time range (top right) is set to "Last 5 minutes" to see the immediate traffic you generated.