

Comprehensive Study Material: Troubleshooting Confluent Kafka

Table of Contents

1. Introduction & Foundational Concepts
2. Broker-Level Troubleshooting
3. Consumer & Producer Troubleshooting
4. Network & Connectivity Issues
5. Security & Authentication Problems
6. Monitoring & Diagnostics
7. Performance & Tuning Issues
8. Common Error Codes & Solutions
9. Best Practices & Preventive Measures

1. Introduction & Foundational Concepts

1.1 Understanding Confluent Kafka Architecture

Confluent Kafka is an enterprise-grade Apache Kafka distribution with additional management tools and features. The core components include:

Brokers: Core Kafka servers that manage partitions and handle produce/consume requests

ZooKeeper/KRaft: Coordination layer (ZooKeeper is legacy; KRaft is the new mode)

Control Center: Web-based monitoring and management interface

Schema Registry: Manages data schemas for structured messages

Kafka Connect: Framework for integrating external systems

ksqlDB: SQL-based stream processing engine

1.2 Troubleshooting Mindset

Effective troubleshooting follows a systematic approach:

1. Identify: Recognize symptoms and anomalies
2. Isolate: Determine affected components (broker, topic, consumer, etc.)
3. Investigate: Analyze logs, metrics, and configurations
4. Remediate: Apply targeted fixes
5. Validate: Verify resolution and prevent recurrence

2. Broker-Level Troubleshooting

2.1 Broker Startup Issues

Problem: Broker fails to start

Common causes and solutions:

- JVM Configuration: Ensure sufficient heap memory is allocated
 - Set `KAFKA_HEAP_OPTS=-Xmx4G -Xms4G` (adjust based on available memory)
 - Default is often too low for production workloads
- Port Conflicts: Verify ports are available
 - Default listener port: 9092
 - Check: `netstat -tuln | grep 9092` (Linux/Mac) or `netstat -ano | findstr 9092` (Windows)
- Log Directory Permissions: Ensure the broker can write to log directories
 - Check: `ls -la /var/kafka/data`

- Fix: `chmod -R 755 /var/kafka/data` and set proper ownership
- Corrupted State Files: ZooKeeper or KRaft state may be corrupted
 - Backup and remove state: `rm -rf /var/kafka/data/_uuid_registry_metadata.log`
 - Restart broker to recreate

Diagnostic Steps:

```
# Check broker logs  
tail -f /var/log/kafka/server.log  
  
# Verify JVM startup  
ps aux | grep kafka
```

```
# Test broker connectivity  
nc -zv localhost 9092
```

2.2 Broker Crashes & Recovery

Problem: Broker crashes unexpectedly

Investigation Checklist:

1. Review Error Logs
 - Look for OutOfMemoryError, SocketException, or IOException
 - Pattern: `grep -i "error\|exception" /var/log/kafka/server.log | tail -50`

2. Memory Issues

- Monitor with: `jps -m` (Java processes with memory)
- JMX metrics: `kafka.server:type=BrokerTopicMetrics,name=*`
 - If heap exhausted: increase `KAFKA_HEAP_OPTS` and `KAFKA_JVM_PERFORMANCE_OPTS`

3. Disk Space

- Kafka cannot write logs if disk is full
- Check: `df -h /var/kafka/data`
- Solution: Delete old segments or expand disk

4. Network Saturation

- Monitor: `netstat -s` for dropped packets
- Check broker's `NetworkMetrics` via JMX
- Adjust: `num.network.threads` and `num.io.threads`

Recovery Procedure:

- Single broker failure: Replication handles failover automatically
- If leader is affected: ISR (In-Sync Replicas) promote next replica
- Full cluster failure: Restart brokers in any order; leader election occurs

2.3 Partition & Replication Issues

Problem: Under-replicated partitions

This is a critical issue indicating data loss risk.

Symptoms:

- Control Center shows "Under Replicated Partitions"
- Metric: `kafka.server:type=ReplicaManager,name=UnderReplicatedPartitions`

Causes & Solutions:

1. Broker Offline

- Check: `bin/kafka-broker-api-versions.sh --bootstrap-server localhost:9092`
- If broker unresponsive: check network, logs, disk space

2. Slow Broker (Replica Lag)

- Metric: `kafka.server:type=ReplicaFetcherManager,name=MaxLag,clientId=Replica`
- Solution: Increase `replica.lag.time.max.ms` or optimize broker resources

3. ISR (In-Sync Replicas) Shrinkage

- Monitor: `kafka.server:type=ReplicaManager,name=ISRShrinks`
- Root cause: Leader detects replica as stale
- Recovery: Fix the slow/offline broker; ISR will expand when replica catches up

Diagnostic Commands:

```
# Describe topic with details
```

```
kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic my-topic
```

```
# Check replica log end offset
```

```
kafka-replica-verification.sh --broker-list localhost:9092 --topic-white-list ".*"
```

```
# Monitor ISR changes  
kafka-consumer-groups.sh --bootstrap-server localhost:9092 --describe --all-groups
```

2.4 Leader Election Problems

Problem: Leader election takes too long or fails

Key Factors:

1. Unclean Leader Election (`unclean.leader.election.enable=false` recommended)

- If enabled, allows data loss
- When no in-sync replicas available, picks newest replica (potentially lost data)
- Keep disabled; fix underlying replication issues instead

2. Stuck Leader Elections

- Cause: ZooKeeper/KRaft coordination delays
- Check ZooKeeper health: `echo stat | nc localhost 2181`
- Monitor: `kafka.server:type=KafkaRequestHandlerPool,name=AvgIdlePercent`

3. Broker Not Recovering State

- Brokers need time to recover state on startup
- Monitor: `kafka.server:type=ReplicaManager,name=LeaderCount`
- Allow adequate startup time before declaring failure

3. Consumer & Producer Troubleshooting

3.1 Producer Issues

Problem: Producer fails to send messages

Common Error Codes:

1. LEADER_NOT_AVAILABLE

- Meaning: No leader for partition
- Solution: Wait for leader election; check broker health
- Code: Implement exponential backoff retry logic

2. REQUEST_TIMED_OUT

- Meaning: Broker didn't respond in time
- Solutions:
 - Increase `request.timeout.ms` (default: 30s)
 - Increase `delivery.timeout.ms` (default: 2min)
 - Increase broker's `num.network.threads`

3. MESSAGE_TOO_LARGE

- Meaning: Message exceeds max size
- Check: `message.max.bytes` (broker) and `max.request.size` (producer)
- Solution: Increase both or compress messages

4. NOT_LEADER_FOR_PARTITION

- Meaning: Partition leader changed
- Solution: Producer automatically updates metadata; may retry failed sends

Producer Configuration Tuning:

```
properties
# Reliability
acks=all # Wait for all replicas; safest but slower
retries=3
max.in.flight.requests.per.connection=1 # Ensures ordering with retries

# Performance
batch.size=16384 # Bytes; larger = better throughput
linger.ms=10 # Wait up to 10ms for batching
compression.type=snappy # or lz4, gzip

# Timeouts
request.timeout.ms=30000
delivery.timeout.ms=120000
```

Debugging:

```
# Enable producer logs
export KAFKA_LOG4J_OPTS="-Dlog4j.configuration=file:log4j.properties"

# Monitor producer metrics via JMX
# Key metrics:
- kafka.producer:type=producer-metrics,client-id=*
- kafka.producer:type=producer-topic-metrics,client-id=*,topic=*
```

3.2 Consumer Lag Issues

Problem: Consumer group is far behind (high lag)

Understanding Consumer Lag:

- Lag: Difference between latest offset (partition end offset) and consumer's committed offset
- Formula: `Max Offset - Consumer Offset`
- Critical metric for monitoring data freshness

Root Causes:

1. Slow Consumer Processing

- Check consumer processing time: `process.duration_ms`
- Monitor: `poll()` frequency and handler performance
- Solution: Optimize application logic or increase consumer instances

2. Consumer Crashes or Rebalancing

- Rebalancing causes pause in message consumption
- During rebalancing: no progress on lag
- Symptom: Lag spikes correlated with consumer logs

3. Topic Message Rate > Consumer Rate

- Producers sending faster than consumer can process
- Solution: Scale consumers horizontally

4. Network Issues

- Broker-to-consumer latency
- Check: `fetch.wait.max.ms` (broker) and `fetch.min.bytes` (consumer)

Consumer Configuration Tuning:

properties

Parallelism

max.poll.records=500 # Default 500; tune based on message size

fetch.min.bytes=1 # Minimum bytes before returning; default 1

fetch.max.wait.ms=500 # Max wait time; default 500ms

Session management

session.timeout.ms=30000 # Default; increase if rebalancing frequently

heartbeat.interval.ms=10000 # Should be 1/3 of session.timeout.ms

Offset commit

enable.auto.commit=true # or false for manual control

auto.commit.interval.ms=5000 # Commit frequency

Performance

max.partition.fetch.bytes=1048576 # Max bytes per partition; default 1MB

connections.max.idle.ms=540000 # Keep-alive

Monitoring Consumer Lag:

Via Consumer Group Command

kafka-consumer-groups.sh --bootstrap-server localhost:9092 \

```
--group my-consumer-group --describe
```

Output columns:

TOPIC, PARTITION, CURRENT-OFFSET, LOG-END-OFFSET, LAG, CONSUMER-ID, HOST, CLIENT-ID

Interpret:

- LAG = 0: Consumer caught up
- LAG > 0: Consumer behind
- CURRENT-OFFSET increases over time = consumer processing

Remediation Steps:

1. Identify bottleneck: Is it consumer or producer?
2. Increase consumers: Add instances up to partition count
3. Optimize processing: Profile application code
4. Monitor rebalancing: Minimize frequency and duration
5. Set up alerting: Alert when lag exceeds threshold

3.3 Consumer Group Rebalancing Issues

Problem: Excessive rebalancing causing service interruption

Rebalancing Triggers:

1. Consumer joins/leaves group
2. Consumer heartbeat missed

3. Partition count changes
4. Manual partition reassignment

Issues During Rebalancing:

- Stop-the-world: No message processing during rebalance
- Impact: Application unavailability, lag growth
- Duration: Typically 5-30 seconds; longer in large clusters

Solutions:

1. Extend Session Timeout

- Increase `session.timeout.ms` (default: 30s) to 60-90s
- Reduces rebalancing from temporary network glitches
- Trade-off: Delayed detection of failed consumers

2. Stable Assignment Strategy

- Use `partition.assignment.strategy=RoundRobin` or `StickyAssignor`
- `CooperativeStickyAssignor`: Minimizes stop-the-world time (newer, recommended)

3. Graceful Shutdown

```
python
# In consumer code
try:
    while True:
        records = consumer.poll(100)
        # process records
except KeyboardInterrupt:
```

```
consumer.close() # Graceful close triggers cooperative rebalance
```

4. Monitor Rebalancing Events

- Metric: `kafka.consumer:type=consumer-coordinator-metrics,client-id=*`
- Look for: `commit-latency-avg` and `assigned-partitions` changes

4. Network & Connectivity Issues

4.1 Connection Failures

Problem: Cannot connect to Kafka cluster

Diagnostic Checklist:

1. Verify Bootstrap Servers

```
# Test DNS resolution  
nslookup kafka-broker-1.example.com  
  
# Test TCP connectivity  
telnet kafka-broker-1.example.com 9092  
or nc -zv kafka-broker-1.example.com 9092
```

2. Firewall Rules

- Ensure ports are open: 9092 (plaintext), 9093 (SSL), 9094 (SASL)
- Check: `sudo iptables -L -n` or cloud provider's security groups

3. Broker Listener Configuration

- In `server.properties`:

```
listeners=PLAINTEXT://0.0.0.0:9092
```

```
advertised.listeners=PLAINTEXT://kafka-broker-1.example.com:9092
```

- `listeners`: Internal binding

- `advertised.listeners`: What clients use to connect back

4. Client Configuration

```
bootstrap.servers=kafka-broker-1.example.com:9092,kafka-broker-2.example.com:9092
```

4.2 Metadata Refresh Issues

Problem: Broker list updates fail; clients stuck with stale metadata

Root Causes:

1. Advertised Listeners Misconfigured

- Broker advertises unreachable address
- Check: `kafka-broker-api-versions.sh --bootstrap-server localhost:9092`
- Verify advertised addresses are reachable from client perspective

2. DNS Inconsistency

- Broker hostname resolves differently on client vs. broker
- Solution: Use IPs or consistent FQDN

3. Metadata Cache Stale

- Client caches metadata; refreshes on errors or timeout
- `metadata.max.age.ms`: 5min default
- Force refresh: Retry on MetadataException

4.3 Inter-Broker Communication Issues

Problem: Brokers cannot communicate with each other

Symptoms:

- Replication lag
- Leader election delays
- Under-replicated partitions

Diagnosis:

1. Check Inter-Broker Protocol

- `inter.broker.protocol.version` in `server.properties`
- All brokers must support same version

2. Test Broker-to-Broker Connectivity

```
# From broker A, test connectivity to broker B
```

```
nc -zv broker-b.example.com 9092
```

```
# Check broker logs for connection errors
```

```
grep "Connection refused\|Network is unreachable" /var/log/kafka/server.log
```

3. Verify Replication Configuration

- `min.insync.replicas`: Should be < replication factor
- `default.replication.factor`: Ensure adequate replicas

5. Security & Authentication Problems

5.1 SSL/TLS Certificate Issues

Problem: SSL handshake failures

Common Errors:

1. CERTIFICATE_VERIFY_FAILED

javax.net.ssl.SSLHandshakeException: sun.security.validator.ValidatorException

- Cause: Client doesn't trust broker's certificate

- Solution:

```
# Import broker certificate to client truststore
```

```
keytool -import -alias kafka-broker -file broker-cert.pem \  
-keystore client-truststore.jks
```

2. HOSTNAME_MISMATCH

- Certificate CN or SAN doesn't match broker hostname
- Solution: Regenerate certificate with correct hostname

3. CERTIFICATE_EXPIRED

- Check: `openssl x509 -in broker-cert.pem -noout -dates`
- Renew certificate before expiration

Client Configuration for SSL:

properties

```
# Producer/Consumer  
security.protocol=SSL  
ssl.truststore.location=/path/to/truststore.jks  
ssl.truststore.password=password  
ssl.keystore.location=/path/to/keystore.jks # For mTLS  
ssl.keystore.password=password  
ssl.key.password=password
```

5.2 SASL Authentication Issues

Problem: SASL authentication fails

Common Scenarios:

1. SASL/PLAIN (Username/Password)

```
properties  
security.protocol=SASL_PLAINTEXT  
sasl.mechanism=PLAIN  
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule \  
required username="user1" password="secret";
```

- Verify credentials in broker's `jaas.config`

2. SASL/SCRAM

- More secure than PLAIN

```
# Create user in broker
kafka-configs.sh --bootstrap-server localhost:9092 \
--entity-type users --entity-name user1 \
--alter --add-config 'SCRAM-SHA-256=[password=secret]'
```

3. SASL/GSSAPI (Kerberos)

- Requires Kerberos infrastructure
- Check: Kerberos ticket validity

```
klist # List current tickets
kinit -R # Renew ticket
```

Diagnosis:

```
# Enable JAAS logging
export KAFKA_DEBUG_OPTS="-Dcom.sun.security.sasl.debug=true"

# Check broker's JAAS configuration
grep "KafkaServer" /path/to/jaas.conf
```

5.3 Authorization (ACL) Issues

Problem: Access denied despite correct authentication

Error: TOPIC_AUTHORIZATION_FAILED

Diagnosis & Fix:

```
# List current ACLs
```

```
kafka-acls.sh --bootstrap-server localhost:9092 --list
```

```
# Grant permissions
```

```
kafka-acls.sh --bootstrap-server localhost:9092 \
--add --allow-principal User:alice \
--operation Read,Write \
--topic my-topic
```

```
# Grant consumer group permissions
```

```
kafka-acls.sh --bootstrap-server localhost:9092 \
--add --allow-principal User:alice \
--operation Read \
--group my-consumer-group
```

```
# Grant cluster admin
```

```
kafka-acls.sh --bootstrap-server localhost:9092 \
--add --allow-principal User:alice \
--operation ClusterAction \
--cluster
```

6. Monitoring & Diagnostics

6.1 JMX Metrics Overview

Key Broker Metrics:

Metric	Purpose	Alert Threshold
`BytesInPerSec` / `BytesOutPerSec`	Throughput	Monitor trends
`FetchConsumerTotalTimeMs`	Consumer fetch latency	> 500ms sustained
`ProduceLocalTimeMsMean`	Producer latency	> 100ms sustained
`UnderReplicatedPartitions`	Data risk	> 0 (critical)
`OfflinePartitionsCount`	Unavailable partitions	> 0 (critical)
`ISRShrinks` / `ISRExpands`	Replica issues	Track frequency
`NetworkProcessorAvgIdlePercent`	Network capacity	< 20% = near limit
`RequestHandlerAvgIdlePercent`	I/O capacity	< 20% = near limit

Key Consumer Metrics:

Metric	Purpose	Alert Threshold
`ConsumeRate`	Messages/sec consumed	Monitor trends
`FetchLatencyAvg`	Time to fetch batch	> 100ms
`CommitLatencyAvg`	Offset commit time	> 50ms
`JoinRate`	Consumer joins/sec	Spike = rebalancing
`SyncRate`	Group syncs/sec	Spike = rebalancing

6.2 Diagnostic Tools

1. Confluent Diagnostics Bundle Tool

```
# Generate comprehensive diagnostic bundle
confluent-diagnostics.sh --broker-list localhost:9092 \
--output-directory /tmp/kafka-diagnostics
```

Bundle includes:

- JMX metrics snapshots
- Broker configurations
- Broker and system logs
- Zookeeper state (if applicable)
- Topic/partition metadata

2. Control Center UI

- Navigate to: `http://localhost:9021`
- Sections:
 - Cluster: Overall health, broker status
 - Topics: Partition distribution, replication
 - Consumer Groups: Lag, throughput
 - Alerts: Predefined issue notifications

3. Command-Line Tools

```
# Topic metadata
kafka-topics.sh --bootstrap-server localhost:9092 --describe

# Consumer group details
kafka-consumer-groups.sh --bootstrap-server localhost:9092 \
--group my-group --describe --members

# Replication verification
kafka-replica-verification.sh --broker-list localhost:9092 \
--topic-white-list ".*"

# Broker API versions
kafka-broker-api-versions.sh --bootstrap-server localhost:9092
```

6.3 Log Analysis

Critical Log Patterns:

1. Broker Startup Issues

[KafkaServer id=1] started successfully -> Broker ready

Fatal error during KafkaServer startup -> Fatal error

ERROR Error starting socket server -> Port/binding issue

2. Replication Problems

ISR for partition [topic-0] shrinks to [1,2] -> Replica lag

Replica 3 is not in ISR for partition [topic-0] -> Replica offline

Leader election took 5234ms -> Slow election

3. Security Issues

[SaslAuthentication] Failed authentication -> SASL failure

NOT_AUTHORIZED -> ACL denial

SSL/TLS: Unsupported or unrecognized SSL message -> SSL config error

7. Performance & Tuning Issues

7.1 Throughput Bottlenecks

Problem: Cluster not achieving expected throughput

Diagnosis Steps:

1. Identify Bottleneck Layer

- Network: Monitor `BytesInPerSec` vs. network capacity
- Disk I/O: Monitor `iowait` and disk latency
- CPU: Monitor `RequestHandlerAvgIdlePercent`
- Memory: Monitor GC frequency and pause time

2. Network Bottleneck Solutions

- Increase `num.network.threads` (default: 3)
- Enable batching: `batch.size=32768` (producer)
- Compress: `compression.type=snappy`

3. Disk I/O Bottleneck Solutions

- Increase `num.io.threads` (default: 8)
- Use faster disks (SSD)
- Tune OS: increase `vm.dirty_ratio` and `vm.dirty_background_ratio`

4. CPU Bottleneck Solutions

- Compression trade-off: CPU vs. network bandwidth
- Profile application code
- Increase broker instances (horizontal scaling)

7.2 Latency Issues

Problem: High end-to-end latency (producer → broker → consumer)

Measurement Points:

Producer Broker Consumer

↓ ↓ ↓

[send]→[enqueue]→[persist]→[fetch]→[consume]

Tuning Strategies:

Component	High Latency Cause	Solution
Producer	Batching wait	Reduce `linger.ms`
Producer	Acks wait	Change to `acks=1` (less safe)
Broker	Replication sync	Reduce `min.insync.replicas` (risky)
Consumer	Polling interval	Increase `max.poll.records`
Consumer	Processing time	Optimize application logic

7.3 Memory Management

Problem: Frequent GC pauses or OutOfMemoryError

JVM Tuning:

```
# In kafka-env.sh or server startup  
  
export KAFKA_HEAP_OPTS="-Xmx4G -Xms4G"  
  
export KAFKA_JVM_PERFORMANCE_OPTS="-XX:+UseG1GC \  
-XX:MaxGCPauseMillis=20 \  
-XX:InitiatingHeapOccupancyPercent=35 \  
-XX:G1HeapRegionSize=16M"
```

Broker-Level Tuning:

- `log.flush.interval.messages`: Reduce to flush pages sooner (default: disabled)
- `log.segment.bytes`: Smaller segments reduce memory pressure (default: 1GB)
- `replica.socket.receive.buffer.bytes`: Default 64KB; reduce if memory-constrained

8. Common Error Codes & Solutions

Error Code Reference

Code	Name	Cause	Solution
0	NONE	Success	N/A
1	OFFSET_OUT_OF_RANGE	Consumer seeks past latest offset	Reset to latest: `kafka-consumer-groups.sh --reset-offsets --to-latest`
2	INVALID_COMMIT_OFFSET	Invalid offset commitment	Check consumer logic; may need offset reset
3	NOT_COORDINATOR_FOR_GROUP	Client queried wrong broker	Client automatically retries with correct coordinator
5	LEADER_NOT_AVAILABLE	No partition leader	Wait for leader election; check broker health
6	NOT_LEADER_FOR_PARTITION	Stale broker in metadata	Producer/consumer auto-update metadata
7	REQUEST_TIMED_OUT	Broker slow/unresponsive	Increase timeouts; check broker resources
10	MESSAGE_TOO_LARGE	Message exceeds max size	Increase `message.max.bytes` (broker) and `max.request.size` (client)
31	POLICY_VIOLATION	Topic creation violates policy	Check `create.topic.policy.class.name`
33	INVALID_PRINCIPAL_TYPE	Unknown principal type in ACL	Use correct format: `User:alice`, `User:*`
40	SASL_AUTHENTICATION_FAILED	Wrong SASL credentials	Verify username/password or Kerberos setup
41	UNSUPPORTED_SASL_MECHANISM	Broker doesn't support mechanism	Configure matching SASL mechanism on client and broker
45	TOPIC_AUTHORIZATION_FAILED	ACL denies operation	Grant required permissions via `kafka-acls.sh`

| 47 | SSL_HANDSHAKE_FAILURE | SSL/TLS error | Verify certificates; check SSL configuration |

9. Best Practices & Preventive Measures

9.1 Operational Best Practices

1. Capacity Planning

- Monitor and forecast disk growth (log retention)
- Size brokers for 70-80% capacity max
- Plan for 3x peak throughput headroom

2. Configuration Management

- Version control all configurations
- Document non-default settings
- Test changes in staging before production

3. Monitoring Strategy

- Set up alerts for critical metrics:
 - UnderReplicatedPartitions > 0
 - OfflinePartitionsCount > 0
 - ConsumerLag > threshold
 - RequestHandlerAvgIdlePercent < 20%
 - DiskUsage > 80%
- Establish runbook for each alert

4. Backup & Disaster Recovery

- Regularly backup ZooKeeper data (or KRaft snapshots)
- Document recovery procedures
- Test recovery in staging environment

9.2 Configuration Best Practices

properties

```
# Broker Configuration (server.properties)
```

Replication & Durability

```
min.insync.replicas=2 # Ensure quorum writes  
default.replication.factor=3 # Tolerate 1 broker failure  
unclean.leader.election.enable=false # Prevent data loss
```

Performance

```
num.network.threads=8 # Adjust based on cores  
num.io.threads=8 # Disk I/O workers  
num.replica.fetchers=4 # Replication parallelism  
compression.type=snappy # Default compression
```

Timeouts & Health

```
replica.lag.time.max.ms=30000 # Max replica lag  
group.min.session.timeout.ms=6000 # Consumer session timeout  
group.max.session.timeout.ms=300000 # Max session timeout
```

Retention

```
log.retention.hours=168 # 7 days
```

```
log.cleanup.policy=delete # or 'compact' for compaction  
log.segment.bytes=1073741824 # 1GB segments  
  
# Monitoring  
auto.create.topics.enable=false # Require explicit topic creation  
allow.unclean.shutdown=false # Ensure clean shutdown
```

9.3 Producer Best Practices

```
properties  
# Reliability  
acks=all # Wait for all replicas  
retries=Integer.MAX_VALUE # Unlimited retries  
max.in.flight.requests.per.connection=1 # Ensure ordering  
  
# Performance  
batch.size=32768 # Batch size in bytes  
linger.ms=100 # Wait time for batching  
compression.type=snappy # Reduce network bandwidth  
  
# Fault Tolerance  
request.timeout.ms=30000 # Request timeout  
delivery.timeout.ms=120000 # Total delivery timeout  
enable.idempotence=true # Prevent duplicates
```

9.4 Consumer Best Practices

properties

Parallelism

```
max.poll.records=500 # Messages per poll  
fetch.min.bytes=1024 # Min fetch size  
fetch.max.wait.ms=500 # Max wait time
```

Session Management

```
session.timeout.ms=30000 # Consumer session timeout  
heartbeat.interval.ms=10000 # Heartbeat frequency  
max.poll.interval.ms=300000 # Max processing time
```

Offset Management

```
enable.auto.commit=true # Auto-commit or manual  
auto.commit.interval.ms=5000 # Commit frequency  
auto.offset.reset=earliest # or 'latest' / 'none'
```

Partition Assignment

```
partition.assignment.strategy=org.apache.kafka.clients.consumer.CooperativeStickyAssignor
```

9.5 Troubleshooting Checklists

When Brokers Are Down:

- Check system logs: `/var/log/syslog`
- Verify disk space: `df -h`
- Check open file limits: `ulimit -n`

- Verify port availability: `netstat -tuln`
- Review broker logs: `tail -200 server.log`
- Check JVM memory: `jps -m`
- Verify network connectivity to ZK/KRaft: `nc -zv zk-host`

When Consumers Are Lagging:

- Check consumer processing time
- Verify consumer instance count vs. partitions
- Monitor broker metrics (capacity)
- Check for rebalancing frequency
- Verify network latency: `ping broker`
- Review consumer logs for exceptions
- Check topic throughput vs. consumer rate

When Replication Is Unhealthy:

- Verify all brokers are up: `kafka-broker-api-versions.sh`
- Check ISR status: `kafka-topics.sh --describe`
- Monitor replica lag: `kafka-replica-validation.sh`
- Check inter-broker connectivity: `nc -zv broker-host`
- Verify broker configuration compatibility
- Review broker logs for network errors
- Check disk space on all brokers

10. Advanced Troubleshooting Scenarios

Scenario 1: Cascading Broker Failures

Situation: Broker 1 goes down; shortly after, brokers 2 and 3 become unresponsive.

Root Cause Analysis:

- Broker 1 fails → Partitions lose leader
- Leader election triggers on brokers 2 & 3
- All brokers simultaneously busy → Network timeout from ZooKeeper
- ZooKeeper considers brokers 2 & 3 dead
- Cascading failure

Resolution:

1. Restart broker 1 first (usually enough)
2. If not, restart brokers sequentially (30-60s between each)
3. Increase ZooKeeper session timeout: `zookeeper.session.timeout.ms=40000`
4. Increase broker's ZooKeeper sync time: `zookeeper.sync.time.ms=10000`

Scenario 2: Producer Deadlock with min.insync.replicas

Situation: Producer sends message; hangs indefinitely.

Root Cause:

- `min.insync.replicas=2` configured
- Only 1 broker currently in ISR (other is slow)
- Producer waits for 2 acknowledgments; never arrives
- Timeout occurs after `request.timeout.ms`

Resolution:

1. Fix underlying broker issue (check slow broker)
2. Temporarily reduce `min.insync.replicas` if critical
3. Set `request.timeout.ms` to finite value
4. Implement client-side timeout handling

Scenario 3: Consumer Group Stuck During Rebalancing

Situation: Consumer group shows `PreparingRebalance` state for hours.

Root Cause:

- Group coordinator unavailable or unresponsive
- Consumer fails during rebalance (throws exception)
- Broker GC pause during assignment computation

Resolution:

1. Check coordinator broker: `kafka-consumer-groups.sh --describe --group my-group`
2. Verify coordinator broker is healthy
3. Restart consumer application
4. Manually trigger rebalance: `kafka-consumer-groups.sh --reset-offsets --to-latest --group my-group`

Conclusion

Troubleshooting Confluent Kafka effectively requires:

1. Understanding Architecture: Know how brokers, replicas, and consumers interact
2. Systematic Approach: Gather facts, isolate issues, apply targeted fixes

3. Monitoring: Proactive metrics and logging prevent crises
4. Documentation: Maintain runbooks and configurations
5. Testing: Validate fixes in staging before production

Key takeaways:

- Replication health is critical; monitor under-replicated partitions
- Consumer lag indicates end-to-end system health
- Configuration matters: Many issues stem from suboptimal settings
- Timeouts are often symptoms, not root causes; dig deeper
- Preventive measures (alerting, backups, capacity planning) are valuable

Keep this guide nearby for quick reference, and build experience through hands-on troubleshooting in test environments.

Additional Resources

- Confluent Documentation: <https://docs.confluent.io/>
- Apache Kafka Documentation: <https://kafka.apache.org/documentation/>
- Kafka Protocol Documentation: <https://kafka.apache.org/protocol>
- Community Forums: <https://www.confluent.io/community/>
- Confluent Support: <https://support.confluent.io/>